

dbt Training · Course Overview

What you'll learn — and why it matters

● Tier 1

Foundations · ~9 h

● Tier 2

Working Effectively · ~6 h

● Tier 3

Production · ~6 h

17 modules · ~21 hours · dbt Core + Snowflake.

By the end you can **read, write, test, and maintain** dbt models confidently.

Course Map

● Tier 1 — Foundations ~9 h

01 What is dbt and Why We Use It

02 Project Setup, Repo Structure & Execution Sequence

03 Jinja Basics for dbt

04 Materializations

05 Sources and the Medallion Architecture

06 Testing Data Quality

07 Documentation

Full lesson plans + exercises

● Tier 2 — Working Effectively ~6 h

08 Materializations in depth

09 Seeds and variables

10 Macros

11 Slowly changing dimensions and snapshots

12 Selectors, tags, and running subsets

Slides + hands-on practice

● Tier 3 — Production & Advanced ~6 h

13 Workflow: One repo, branching, PRs

14 Advanced testing patterns

15 Incremental patterns for large tables

16 Custom macros and `scd2_merge` deep dive

17 Governance, contracts, and access

Planned — builds on Tier 2

Each module builds on the previous one. Practice will create one complete project, step by step.

layout: default

background: '#f9f8f5'

highlighter: shiki

Before we start · Quick check-in

What do you already know?

No wrong answers — this helps calibrate where we start.

QUESTION 01

Why does dbt exist?

What pain was it built to solve?

QUESTION 02

Where in the data stack does dbt live?

Ingestion? Transformation? Serving?

QUESTION 03

Programming language, IDE, or framework?

How would you describe it to a colleague?

QUESTION 04

What does dbt actually produce?

Tables, dashboards, pipelines, reports?

QUESTION 05

What's the biggest win over raw SQL?

You could write transforms without dbt — so why bother?

QUESTION 06

Analyst tool, engineer tool, or both?

Who writes and runs dbt models in practice?

Interactive Session: Shout out answers. We'll cover all of these in Module 01.